

Prakhar Srivastava

✉ srivastavaprakhar1@gmail.com | in/prakharcode | github/prakharcode | prakhar.sh

EXPERIENCE

Data and AI Engineer, Catawiki — Amsterdam, NL

Apr 2026 – Present

- Built a fully autonomous **agent orchestration** system to triage and resolve operational issues, reducing operational incidents by **60%** via autonomous bot agents.
- Developing a **natural language to SQL** internal bot, enabling business teams to run queries independently on **BigQuery**.
- Reduced pipeline costs by **30%** while optimizing jobs to run faster.

Senior Data Engineer, Tiqets — Amsterdam, NL

Jun 2023 – Mar 2026

- Owning the near real-time EL of data from production systems to the data warehouse, setting up multiple streaming (micro-batch) jobs on Databricks.
- Helped model the customer data platform tables to generate business reporting based on medallion architecture using **pyspark** and **dbt**.
- Owned the migration of complete Airflow infrastructure. Migrated from 2.7.3 to 3.0.6 without any downtime to critical business processes.
- Migrated the data assets to **unity catalog** from Hive Metastore and created a **data governance** framework with organization.
- Helped setting up internal user analysis pipeline to unify mobile and web customer tracking, powering A/B experiments and recommendation pipelines.
- Setting up **MLFlow** on production **Kubernetes** clusters to help ML experimentation and model tracking.
- Moved some of the core ETL pipelines away from **Fivetran** to open-source alternatives like **Meltano** and helped save over **\$4k per month**.
- Helped the team to develop faster with creation of some generic operators and hooks for **Airflow**. Setting up linters like **ruff** and migrating deployments to **uv**.

Data Engineer, TicketSwap — Amsterdam, NL

Sep 2021 – May 2023

- Responsible for creating TicketSwap's entire **data infrastructure** which helped power Business Analytics and Machine Learning pipelines.
- Setting up pipelines to support real-time analytics over **AWS Redshift**. Optimizing **Looker** and **Redshift** to support daily reporting.
- Setting up infrastructure on **AWS** and setting up **Airflow** for scheduled and Ad-Hoc ETL. Creating a data lake using **S3** and **Glue**, later migrated to **Redshift**.
- Brought in **DBT** and together with team created facts and dims for the company.
- Setting up initial flow for ML projects to support ML experiment tracking and development using **MLFlow** and **DBT**.

Data Engineer - II, Atlan — New Delhi / Remote

Nov 2020 – Aug 2021

- Owned **Data Lineage**, a core feature in the Atlan product offering.
- Actively developed and maintained a set of tools that would help in determining the data flow (lineage) inside the organization, independent of the source and destination of the data flow.
- Created an internal framework responsible for scraping information from any source, creating a lineage graph and pushing it to **Apache Atlas DB**. This framework is scalable to any source and any SQL dialect.

Data & Machine Learning Engineer, Atlan – Grid — New Delhi

May 2019 – Oct 2020

- Lead engineering front for **orchestration** and **OLAP/OLTP** engine behind Atlan's product offering.
- Created a multi-stack Presto architecture orchestrated over **Airflow**. To monitor multiple stacks of Presto, created a Presto plugin that sends quantitative and qualitative stats over to **Prometheus** via **Statsd**.
- After applying **multi-stack Presto** architecture, customers realized a cost-benefit of over **\$4k per month**, meeting customer SLAs and bringing processing time from 6 hours to 2 hours for about **15 TB data/day**.

- Part of Atlan's Grid team, actively creating multiple qualitative indicators (such as affluence of a region) using Earth's satellite imagery and open government data.
- Fetched satellite images from *Sentinel-2A* and *Landsat-8* and converted these images into actionable insights, i.e. classification of man-made cover from natural resources like forests, river bodies, etc.
- Using this pipeline we processed around **13 TB** of satellite image data of the entire Indian peninsula. The team created **3.6 Million** data points from which internal indicators were created.

EDUCATION

Guru Gobind Singh Indraprastha University — *New Delhi*

May 2019

B. Tech. Computer Science

- **CGPA 8.15/10.0** — First Division

TALKS

Python and Data Engineering, *PyCon Italy* — *Florence, Italy* 

Jun 2022

- Talked about the Python ecosystem of data engineering — batch pipelines and real-time pipelines using AWS and open-source tools.

Deploy Machine Learning models using Docker, *GDG Delhi – Cloud Meet* — *New Delhi, India* 

Oct 2019

- Covered basics of Docker and modular architecture in ML programming, and creating CI/CD over GitHub for these complex pipelines.

PROJECTS AND CONTRIBUTIONS

Airflow *GitHub://pr/airflow*

- Contributed to the Airflow helm chart, as a bug fix for their user pod creation.

ETL-airflow *GitHub://ETL-airflow*

- Using Airflow to extract weather data from a public source and implementing SQL transformations using Postgres and Airflow.

Large File Process *GitHub://Large-file-processor*

- Using Celery to extract and load data in a parallel batch fashion to a sink database (Postgres).

SKILLS

DE tools/stack: Databricks (PySpark), Airflow, DBT, Athena/Trino, Prometheus, Grafana, Looker

Databases: PostgreSQL, MySQL, DocumentDB, MongoDB

Languages: Python, SQL

Cloud: AWS, Helm, Kubernetes, Terraform

OSS: *Airflow* and Singer SDK